



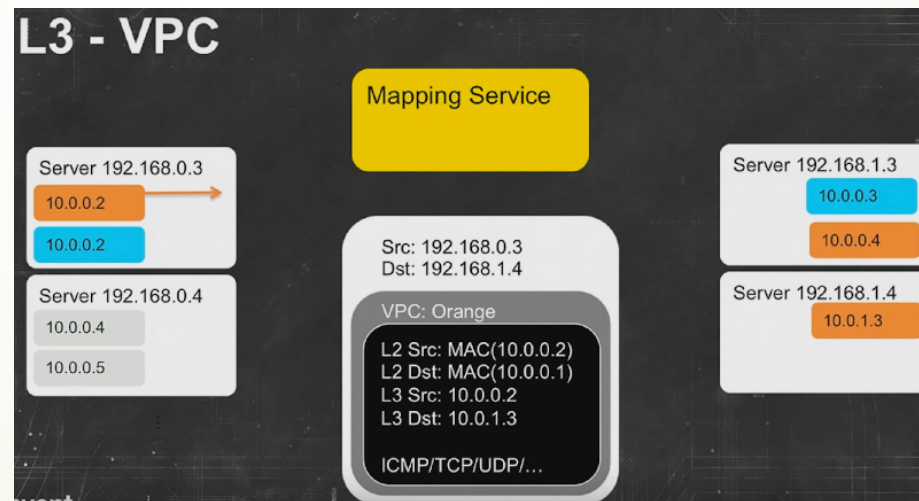
# P4

## Virtual Private Cloud

Based on Amazon AWS VPC presentation  
<https://www.youtube.com/watch?v=Zd5hsL-JNY4>  
<https://www.youtube.com/watch?v=3qln2u1Vr2E>

# Amazon Virtual Private Cloud

- Provide multitenancy in the same hardware infrastructure
- AWS has developed their own encapsulation mechanism
- Centralized mapping service with distributed cache on each compute.
- Each compute cache should have 100% of their required rules.



# Amazon Virtual Private Cloud

## L3 - VPC

Mapping Service

Server 192.168.0.3

10.0.0.2

10.0.0.2

Server 192.168.0.4

10.0.0.4

10.0.0.5

Src: 192.168.0.3  
Dst: 192.168.1.4

VPC: Orange

L2 Src: MAC(10.0.0.2)

L2 Dst: MAC(10.0.0.1)

L3 Src: 10.0.0.2

L3 Dst: 10.0.1.3

ICMP/TCP/UDP/...

Server 192.168.1.3

10.0.0.3

10.0.0.4

Server 192.168.1.4

10.0.1.3



# P4

## Protocol-Independent Packet Processor

- Header
  - definition of the packet structures
- Parsers
  - Parse the packet based on known headers and triggers the entry p
- Control
  - Contains the logic of which tables to be executed
- Tables
  - Match and apply actions
- Actions
  - Packet manipulation



# P4 Private Virtual Cloud

- ARP requests are only replied by P4 Switches
  - Capture ARP request
  - Convert the ARP request into a ARP reply
- IP packets are encapsulated with a custom and non standard protocol called VPC
  - Ethernet type 0x0777
  - Header: Customer id, src/dst IP , src/dst P4 switch
- VPC packets are forwarded by intermediate switches using dst P4 switch label
- Egress switches will remove VPC encapsulation and deliver the IP packet to the node
  - If hosts belong to different networks, src/dst mac addresses will be overwritten
- Non ARP or IP packets are dropped
- ARP or IP packets that does not meet the criteria will be dropped too. For example, an unknown hosts

## L2 – Ethernet (standard)

Compute

10.0.0.2

L2 Ethernet switch

Compute

10.0.0.4



## L2 – Ethernet (standard)

Compute

10.0.0.2



L2 Ethernet switch

Compute

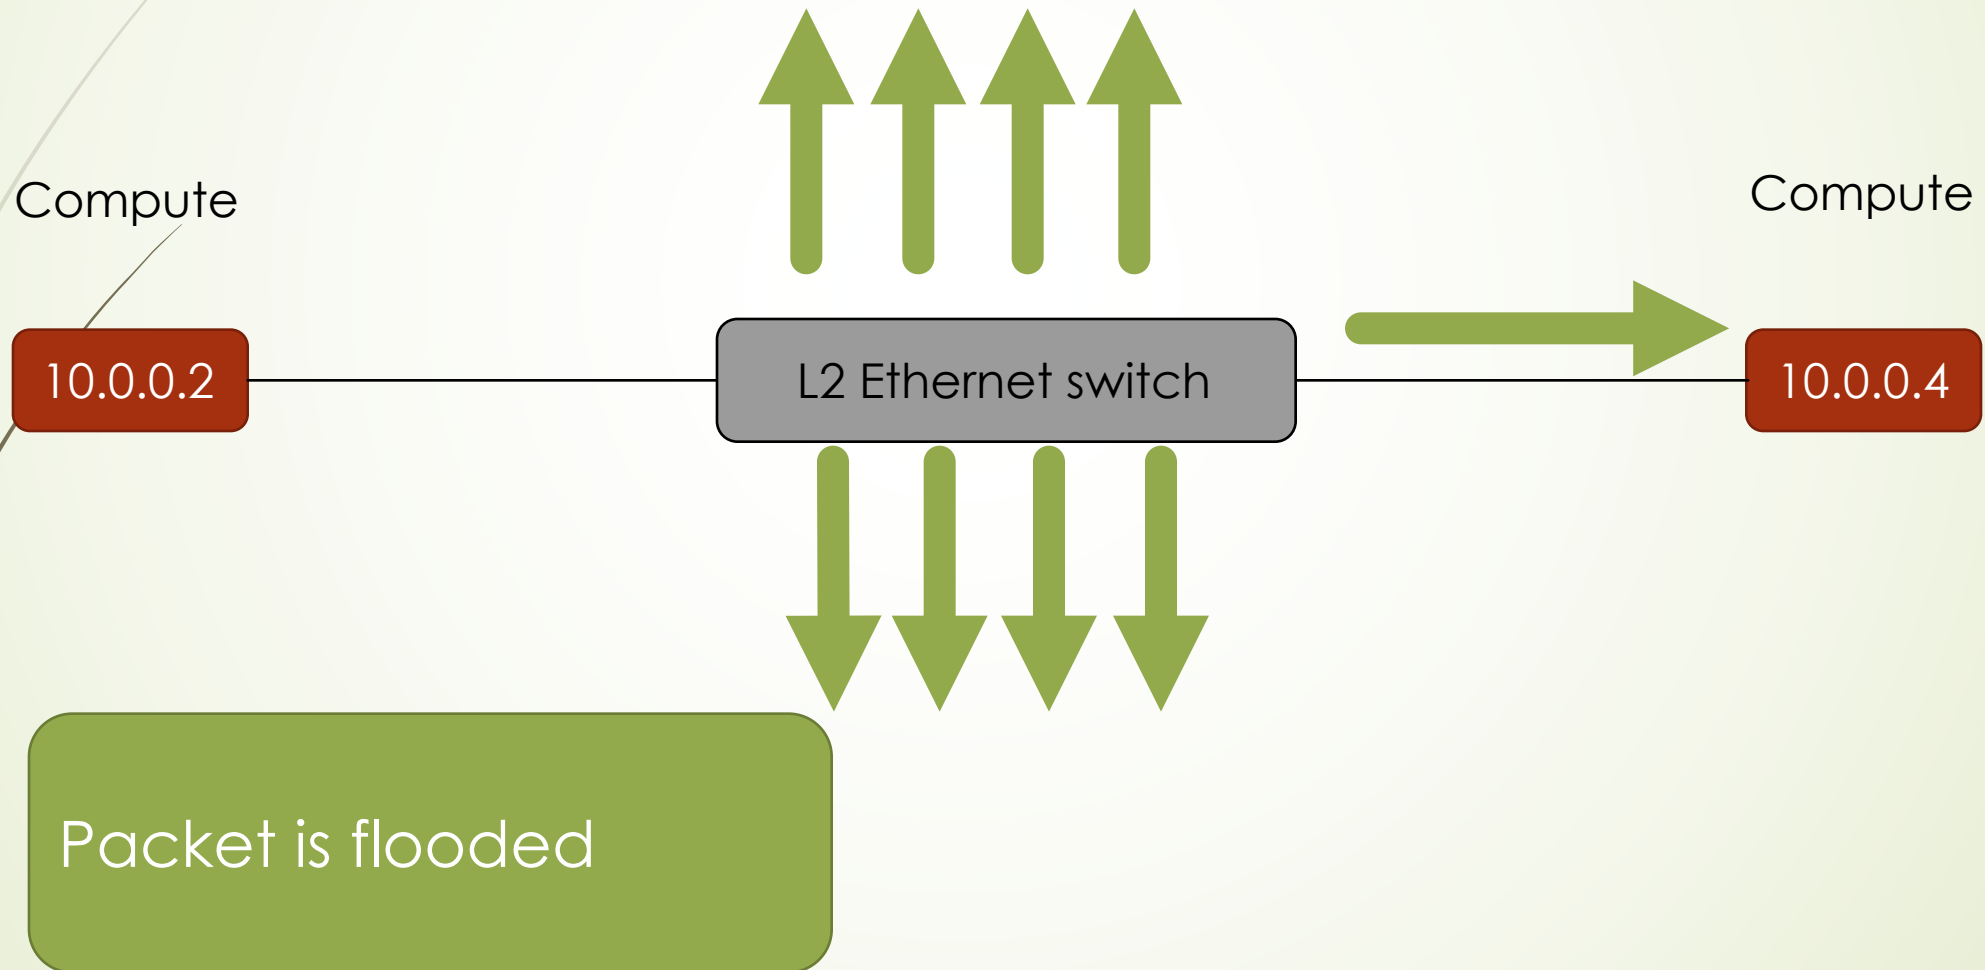
10.0.0.4

L2Src: MAC(10.0.0.2)  
L2Dest: FF:FF:FF:FF:FF:FF

ARP who has 10.0.0.4



## L2 – Ethernet (standard)





## L2 – Ethernet (standard)

Compute

10.0.0.2

L2 Ethernet switch

Compute

10.0.0.4

L2Src: MAC(10.0.0.4)  
L2Dest: MAC(10.0.0.2)

ARP I have 10.0.0.4

## L2 – Ethernet (standard)

Compute

10.0.0.2



L2 Ethernet switch

Compute

10.0.0.4

L2Src: MAC(10.0.0.4)  
L2Dest: MAC(10.0.0.2)

ARP I have 10.0.0.4

## L2 – Ethernet (standard)

Compute

10.0.0.2



L2 Ethernet switch

Compute

10.0.0.4

L2Src: MAC(10.0.0.2)  
L2Dest: MAC(10.0.0.4)

ICMP answer to 10.0.2

## L2 – Ethernet (standard)

Compute

10.0.0.2

L2 Ethernet switch

Compute

10.0.0.4

L2Src: MAC(10.0.0.2)  
L2Dest: MAC(10.0.0.4)

ICMP answer to 10.0.2

## L2 – Ethernet (standard)

Compute

10.0.0.2

L2 Ethernet switch

Compute

10.0.0.4


L2Src: MAC(10.0.0.4)  
L2Dest: MAC(10.0.0.2)

ICMP      reply      from  
10.0.0.4

## L2 – Ethernet (standard)

Compute

10.0.0.2



L2 Ethernet switch

Compute




10.0.0.4

L2Src: MAC(10.0.0.4)  
L2Dest: MAC(10.0.0.2)


ICMP      reply      from  
10.0.0.4

## L2 – Ethernet (using P4 VPC)

Compute



Compute



L2Src: MAC(10.0.0.2)  
L2Dest: FF:FF:FF:FF:FF:FF

ARP who has 10.0.0.4



## L2 – Ethernet (using P4 VPC)

Compute

10.0.0.2

10.0.0.3

10.0.0.2

L2

L2

ARP

```
parser start {  
    return parse_ethernet;  
}  
  
parser parse_ethernet {  
    extract(ethernet);  
    return select(latest.etherType) {  
        ETHERTYPE_VPC : parse_vpc;  
        ETHERTYPE_IPV4 : parse_ipv4;  
        ETHERTYPE_ARP : parse_arp_rarp;  
        default: ingress;  
    }  
}
```

Compute

10.0.0.4

10.0.0.4

10.0.0.5

## L2 – Ethernet (using P4 VPC)

Compute

10.0.0.4

10.0.0.4

10.0.0.5

```
control ingress {
```

```
  if (ethernet.etherType == ETHERTYPE_ARP and arp_rarp.opcode == 1) {  
    apply(address_arp_packet);
```

```
  } else if (ethernet.etherType == ETHERTYPE_IPV4 ){  
    apply(address_ip_packet);
```

```
  }
```

```
  if ((ethernet.etherType == ETHERTYPE_ARP and arp_rarp.opcode == 1) or (ethernet.etherType == ETHERTYPE_IPV4 )){  
    apply(vpc_customer);
```

```
  }
```

```
  if (ethernet.etherType == ETHERTYPE_ARP and arp_rarp.opcode == 1) {  
    if (ingress_metadata.customer > 0){  
      apply(arp_reply);  
    }
```

```
  } else if (ethernet.etherType == ETHERTYPE_IPV4 ){  
    if (ingress_metadata.customer > 0){  
      apply(encapsulate_vpc);  
      apply(vpc_sw_id);  
      apply(vpc_dst);  
    }
```

```
  }
```

Compute

10.0.0.4

10.0.0.4

10.0.0.5

ARP who has 10.0.0.4

## L2 – Ethernet (using P4 VPC)

```
action set_address_arp_packet() {  
    modify_field(ingress_metadata.customer, 0);  
    modify_field(ingress_metadata.srcAddr, arp_rarp_ipv4.srcProtoAddr);  
    modify_field(ingress_metadata.dstAddr, arp_rarp_ipv4.dstProtoAddr);  
}  
  
table address_arp_packet {  
    actions {  
        set_address_arp_packet;  
    }  
    size : 1;  
}
```

## L2 – Ethernet (using P4 VPC)

Compute

10.0.0.2

10.0.0.3

10.0.0.2

```
control ingress {  
  
    if (ethernet.etherType == ETHERTYPE_ARP and arp_rarp.opcode == 1) {  
        apply(address_arp_packet);  
    } else if (ethernet.etherType == ETHERTYPE_IPV4 ){  
        apply(address_ip_packet);  
    }  
    if ((ethernet.etherType == ETHERTYPE_ARP and arp_rarp.opcode == 1) or (ether  
        apply(vpc_customer);  
    }  
    if (ethernet.etherType == ETHERTYPE_ARP and arp_rarp.opcode == 1) {  
        if (ingress_metadata.customer > 0){  
            apply(arp_reply);  
        }  
    } else if (ethernet.etherType == ETHERTYPE_IPV4 ){  
        if (ingress_metadata.customer > 0){  
            apply(encapsulate_vpc);  
            apply(vpc_sw_id);  
            apply(vpc_dst);  
        }  
    }  
}
```

Compute

10.0.0.4

10.0.0.4

10.0.0.5

## L2 – Ethernet (using P4 VPC)

```
Com  
    action set_vpc_customer(customer) {  
        modify_field(ingress_metadata.customer, customer);  
    }
```

```
10.0  
10.0  
10.0  
    table vpc_customer {  
        reads {  
            ethernet.srcAddr : exact;  
            ingress_metadata.srcAddr : exact;  
        }  
        actions {  
            _drop;  
            set_vpc_customer;  
        }  
        size : 1024;  
    }
```

oute

0.0.4

0.0.4

0.0.5

## L2 – Ethernet (using P4 VPC)

```
action set_vpc_customer(customer) {
```

```
table_add l2_addr _noop 20000 1111 10.0.0.0/24 10.0.0.3 =>  
table_add l2_addr set_l2_addr 10000 1111 0.0.0.0/0 10.0.0.2 => 00:00:00:01:00  
table_add l2_addr set_l2_addr 10000 1111 0.0.0.0/0 10.0.0.3 => 00:00:00:01:00  
table_add l2_addr set_l2_addr 20000 1111 0.0.0.0/0 10.0.0.2 => 00:00:00:01:00  
table_add l2_addr set_l2_addr 20000 1111 0.0.0.0/0 10.0.0.3 => 00:00:00:01:00  
table_add routing_pvc route_vpc 2222 => 5  
table_add routing_pvc route_vpc 3333 => 5  
table_add vpc_customer set_vpc_customer 00:00:00:00:00:66 10.0.0.2 => 10000  
table_add vpc_customer set_vpc_customer 00:00:00:00:00:67 10.0.0.3 => 10000  
table_add vpc_customer set_vpc_customer 00:00:00:00:00:ce 10.0.0.2 => 20000  
table_add vpc_customer set_vpc_customer 00:00:00:00:00:cf 10.0.0.3 => 20000
```

```
size : 1024;
```

```
}
```



## L2 – Ethernet (using P4 VPC)

Comput

10.0.0.2

10.0.0.3

10.0.0.2

pute

0.0.4

0.0.4

0.0.5

```
control ingress {  
  
    if (ethernet.etherType == ETHERTYPE_ARP and arp_rarp.opcode == 1) {  
        apply(address_arp_packet);  
    } else if (ethernet.etherType == ETHERTYPE_IPV4 ){  
        apply(address_ip_packet);  
    }  
    if ((ethernet.etherType == ETHERTYPE_ARP and arp_rarp.opcode == 1) or (ethernet.etherType == ETHERTYPE_IPV4 )){  
        apply(vpc_customer);  
    }  
    if (ethernet.etherType == ETHERTYPE_ARP and arp_rarp.opcode == 1) {  
        if (ingress_metadata.customer > 0){  
            apply(arp_reply);  
        }  
    } else if (ethernet.etherType == ETHERTYPE_IPV4 ){  
        if (ingress_metadata.customer > 0){  
            apply(encapsulate_vpc);  
            apply(vpc_sw_id);  
            apply(vpc_dst);  
        }  
    }  
}
```



## 1.2 Ethernet (using P4 VPC)

Com

10.0

10.0

10.0

pute

0.0.4


0.0.4

0.0.5

```
action set_arp_reply(hwAddr) {
    modify_field(ethernet.dstAddr, ethernet.srcAddr);
    modify_field(ethernet.srcAddr, hwAddr);
    modify_field(arp_rarp.opcode, 2);
    modify_field(arp_rarp_ipv4.dstHwAddr, arp_rarp_ipv4.srcHwAddr);
    modify_field(arp_rarp_ipv4.dstProtoAddr, arp_rarp_ipv4.srcProtoAddr);
    modify_field(arp_rarp_ipv4.srcHwAddr, hwAddr);
    modify_field(arp_rarp_ipv4.srcProtoAddr, ingress_metadata.dstAddr);
    modify_field(standard_metadata.egress_spec, standard_metadata.ingress_port);
}

table arp_reply {
    reads {
        ingress_metadata.customer : exact;
        ingress_metadata.srcAddr : lpm;
        ingress_metadata.dstAddr : exact;
    }
    actions {
        _drop;
        set_arp_reply;
    }
    size : 1024;
}
```

## L2 – Ethernet (using P4 VPC)



L2Src: MAC(10.0.0.4)  
L2Dest: MAC(10.0.0.2)

ARP I have 10.0.0.4

## L2 – Ethernet (using P4 VPC)

Compute

10.0.0.2

Compute

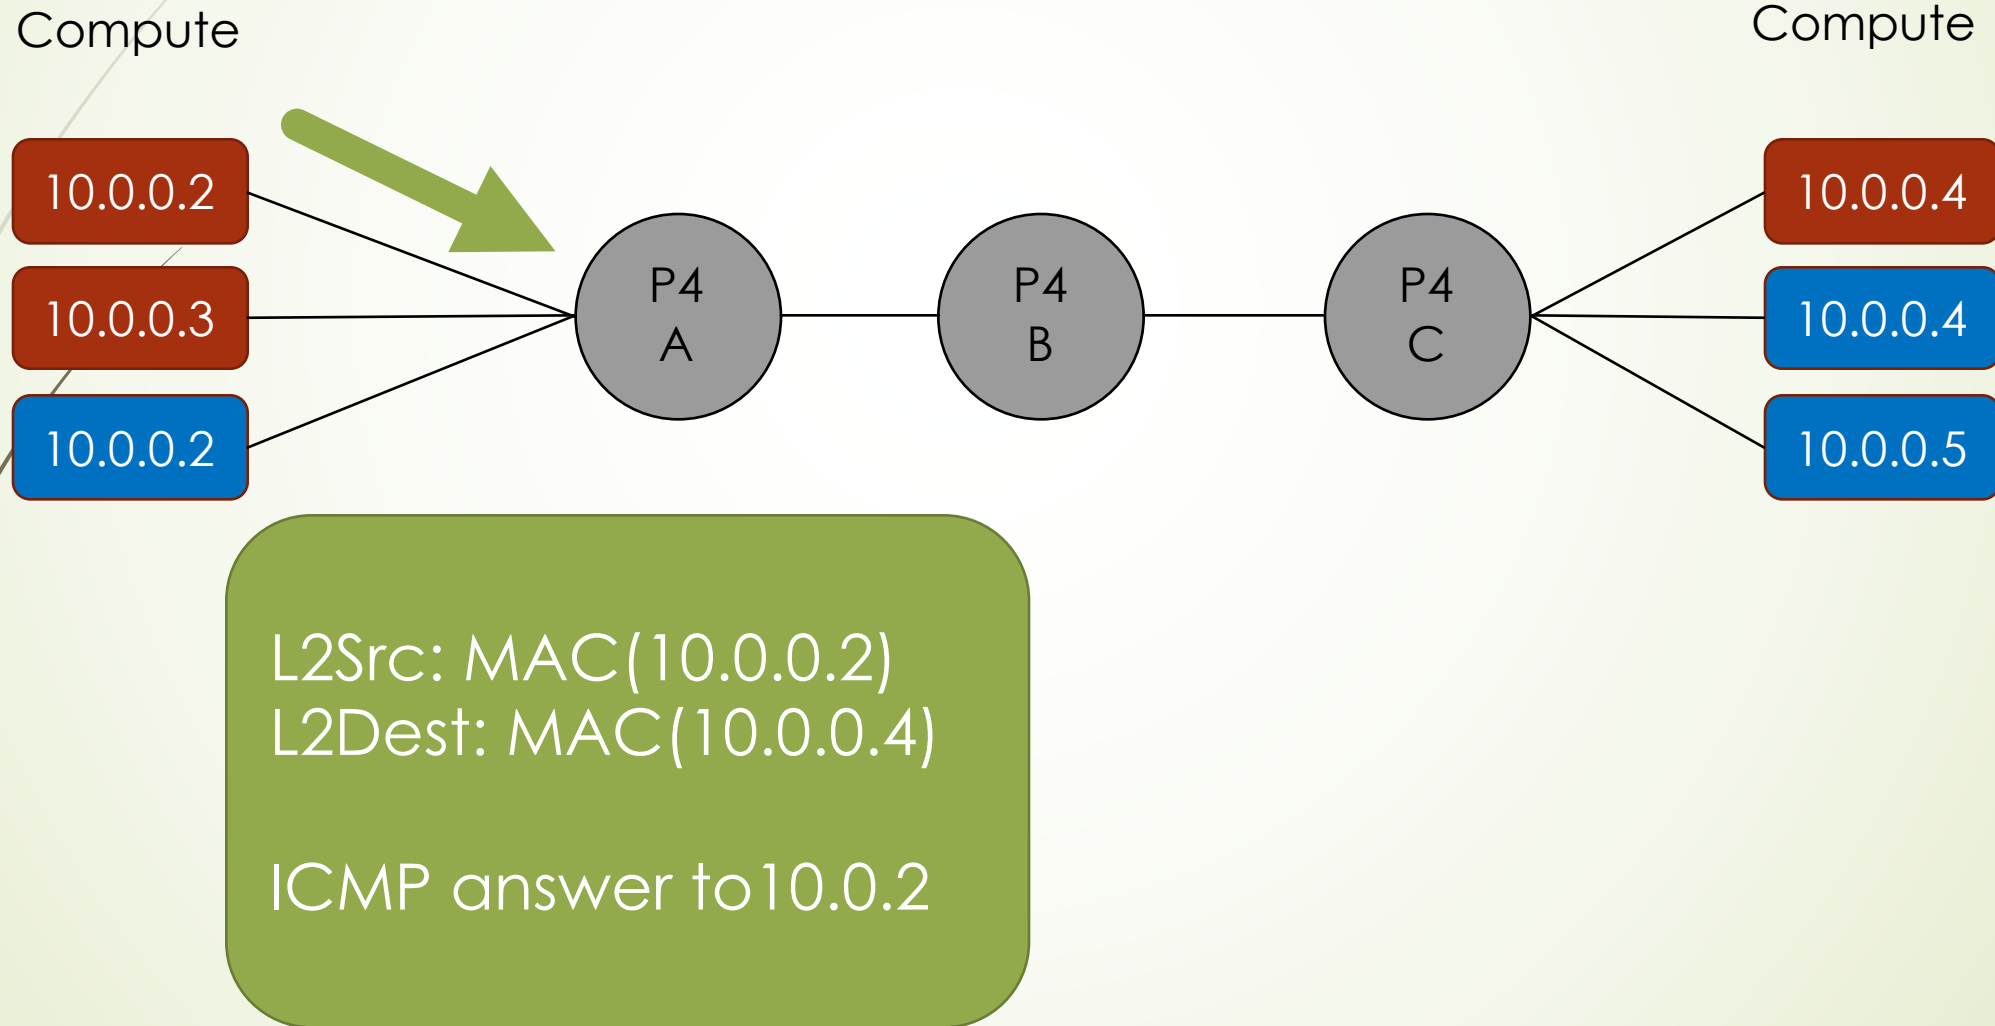
10.0.0.4

**Notice ARP request did not flood and P4 has converted ARP request into ARP reply**

L2Dest: MAC(10.0.0.2)

ARP I have 10.0.0.4

## L2 – Ethernet (using P4 VPC)



## L2 – Ethernet (using P4 VPC)

Compute

10.0.0.2

10.0.0.3

10.0.0.2

```
parser start {  
    return parse_ethernet;  
}  
  
parser parse_ethernet {  
    extract(ethernet);  
    return select(latest.etherType) {  
        ETHERTYPE_VPC : parse_vpc;  
        ETHERTYPE_IPV4 : parse_ipv4;  
        ETHERTYPE_ARP : parse_arp_rarp;  
        default: ingress;  
    }  
}
```

Compute

10.0.0.4

10.0.0.4

10.0.0.5

## L2 – Ethernet (using P4 VPC)

Compute

10.0.0.1

10.0.0.2

10.0.0.3

```
control ingress {  
  
    if (ethernet.etherType == ETHERTYPE_ARP and arp_rarp.opcode == 1) {  
        apply(address_arp_packet);  
    } else if (ethernet.etherType == ETHERTYPE_IPV4 ){  
        apply(address_ip_packet);  
    }  
  
    if ((ethernet.etherType == ETHERTYPE_ARP and arp_rarp.opcode == 1) or (ethernet.etherType == ETHERTYPE_ICMP and icmp.opcode == 0)) {  
        apply(vpc_customer);  
    }  
  
    if (ethernet.etherType == ETHERTYPE_ARP and arp_rarp.opcode == 1) {  
        if (ingress_metadata.customer > 0){  
            apply(arp_reply);  
        }  
    }  
  
    } else if (ethernet.etherType == ETHERTYPE_IPV4 ){  
        if (ingress_metadata.customer > 0){  
            apply(encapsulate_vpc);  
            apply(vpc_sw_id);  
            apply(vpc_dst);  
        }  
    }  
}
```

Compute

10.0.0.4

10.0.0.4

10.0.0.5

ICMP answer to 10.0.2

control ingress {

```
    if (ethernet.etherType == ETHERTYPE_ARP and arp_rarp.opco
        apply(address_arp_packet);
    } else if (ethernet.etherType == ETHERTYPE_IPV4 ){
        apply(address_ip_packet);
    }
    if ((ethernet.etherType == ETHERTYPE_ARP and arp_rarp.opco
        apply(vpc_customer);
    }
    if (ethernet.etherType == ETHERTYPE_ARP and arp_rarp.opco
        if (ingress_metadata.customer > 0){
            apply(arp_reply);
        }
    } else if (ethernet.etherType == ETHERTYPE_IPV4 ){
        if (ingress_metadata.customer > 0){
            apply(encapsulate_vpc);
            apply(vpc_sw_id);
            apply(vpc_dst);
        }
    }
    if (valid(vpc)){
        apply(l2_addr);
        apply(routing_pvc);
        apply(deliver_pvc);
    }
}
```

Compute

10.0.0.2

10.0.0.3

10.0.0.2

Compute

10.0.0.4

10.0.0.4

10.0.0.5



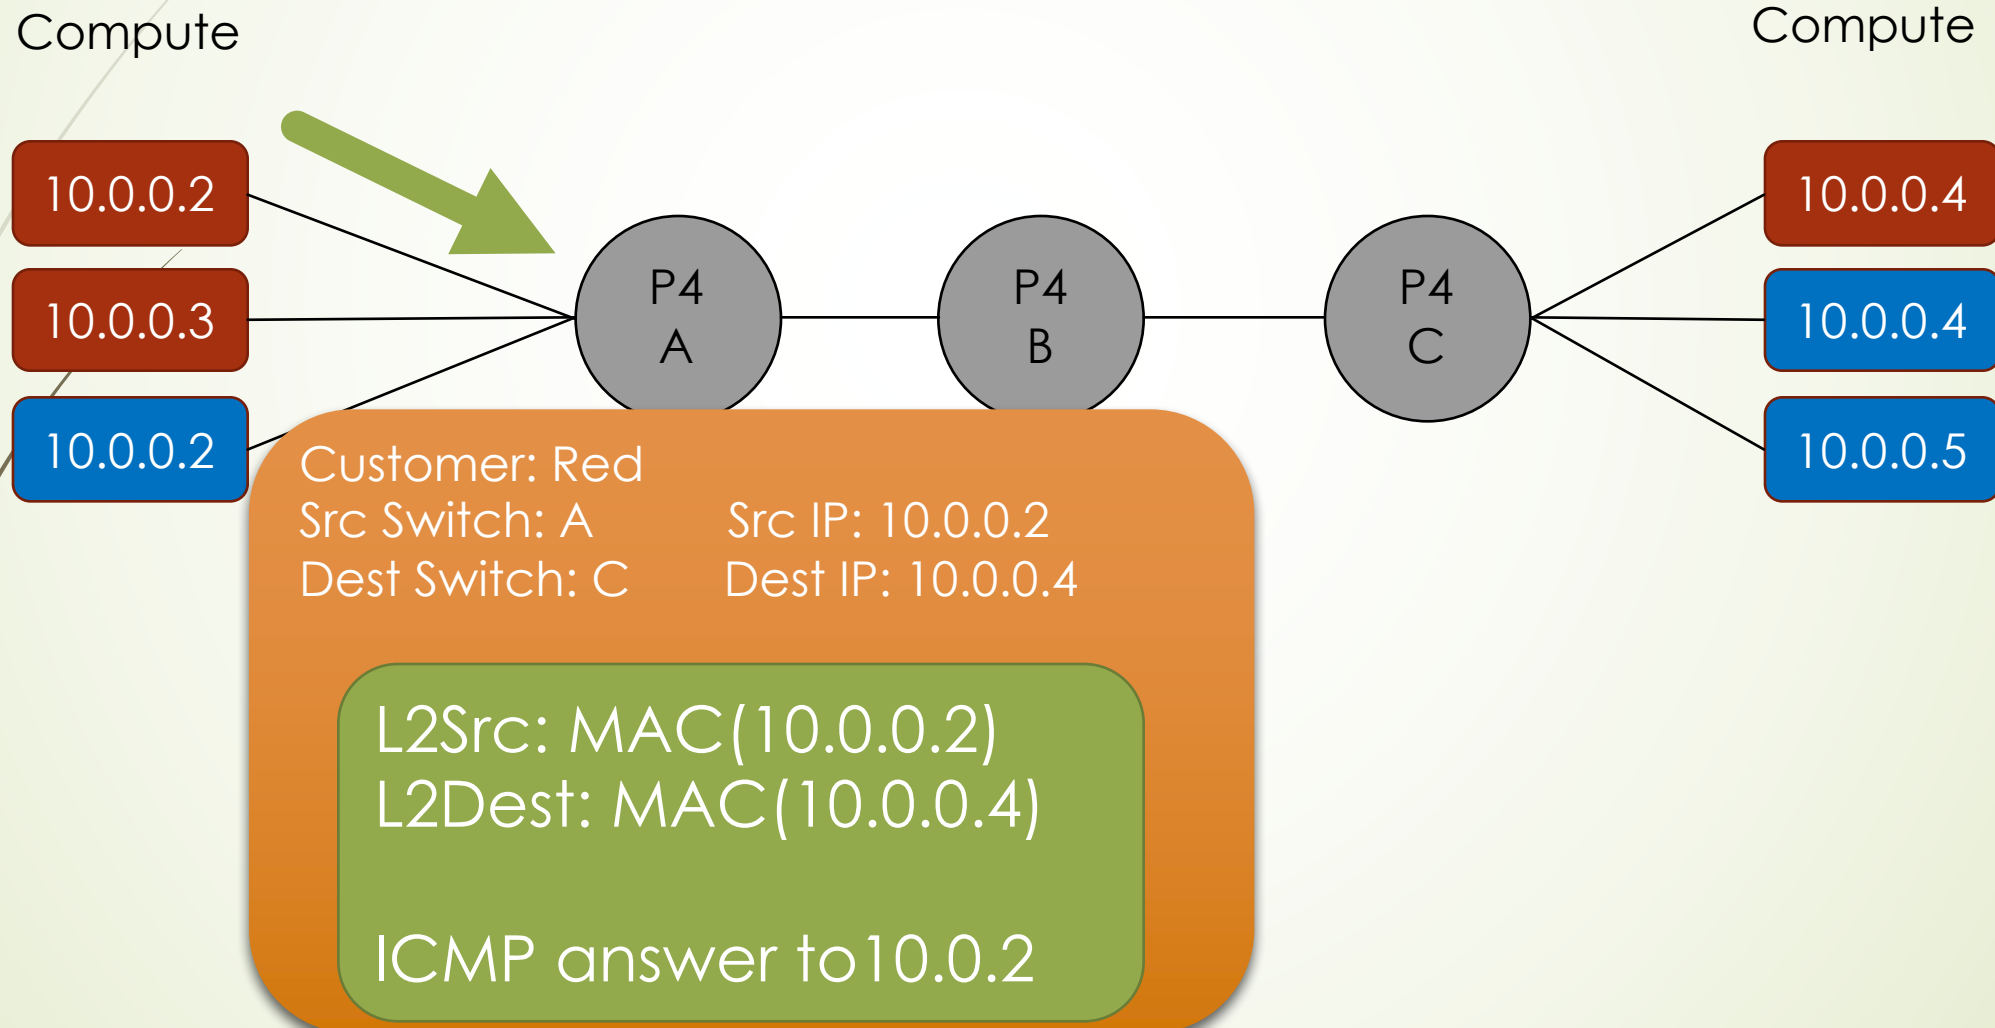
## 10. Firewall (version D4)(DC)

```
action route vpc(port) {  
    modify_field(standard_metadata.egress_spec, port);  
}  
  
table routing_pvc {  
    reads {  
        vpc.dstSw : exact;  
    }  
    actions {  
        route_vpc;  
    }  
    size : 1024;  
}
```

## 10. Firewall (version D4)(D5)




```
action route vpc(port) {  
    modify_field(standard_metadata.egress_spec, port);  
    table_add l2_addr _noop 20000 1111 10.0.0.0/24 10.0.0.3 =>  
    table_add l2_addr set_l2_addr 10000 1111 0.0.0.0/0 10.0.0.2 => 00:00:  
    table_add l2_addr set_l2_addr 10000 1111 0.0.0.0/0 10.0.0.3 => 00:00:  
    table_add l2_addr set_l2_addr 20000 1111 0.0.0.0/0 10.0.0.2 => 00:00:  
    table_add l2_addr set_l2_addr 20000 1111 0.0.0.0/0 10.0.0.3 => 00:00:  
    table_add routing_pvc route_vpc 2222 => 5  
    table_add routing_pvc route_vpc 3333 => 5  
    table_add vpc_customer set_vpc_customer 00:00:00:00:00:00:66 10.0.0.2 =  
    table_add vpc_customer set_vpc_customer 00:00:00:00:00:00:67 10.0.0.3 =  
}  
size : 1024;  
}
```

## L2 – Ethernet (using P4 VPC)




## L2 – Ethernet (using P4 VPC)

Compute



Compute



Customer: Red

Src Switch: A

Dest Switch: C

Src IP: 10.0.0.2

Dest IP: 10.0.0.4




L2Src: MAC(10.0.0.2)

L2Dest: MAC(10.0.0.4)


ICMP answer to 10.0.2

## L2 – Ethernet (using P4 VPC)

Compute



Compute



Customer: Red

Src Switch: A

Dest Switch: C

Src IP: 10.0.0.2

Dest IP: 10.0.0.4

L2Src: MAC(10.0.0.2)

L2Dest: MAC(10.0.0.4)

ICMP answer to 10.0.2

`control ingress {`

```
    if (ethernet.etherType == ETHERTYPE_ARP and arp_rarp.opco
        apply(address_arp_packet);
    } else if (ethernet.etherType == ETHERTYPE_IPV4 ){
        apply(address_ip_packet);
    }
    if ((ethernet.etherType == ETHERTYPE_ARP and arp_rarp.opco
        apply(vpc_customer);
    }
    if (ethernet.etherType == ETHERTYPE_ARP and arp_rarp.opco
        if (ingress_metadata.customer > 0){
            apply(arp_reply);
        }
    } else if (ethernet.etherType == ETHERTYPE_IPV4 ){
        if (ingress_metadata.customer > 0){
            apply(encapsulate_vpc);
            apply(vpc_sw_id);
            apply(vpc_dst);
        }
    }
    if (valid(vpc)){
        apply(l2_addr);
        apply(routing_pvc);
        apply(deliver_pvc);
    }
}
```

Compute

10.0.0.2

10.0.0.3


10.0.0.2

Compute

10.0.0.4

10.0.0.4

10.0.0.5



```
action pop_route_vpc(port) {  
    modify_field(standard_metadata.egress_spec, port);  
    modify_field(ethernet.etherType, vpc.etherType);  
    remove_header(vpc);  
}
```

Control

10.0.0.2

10.0.0.4

10.0.0.5

```
table deliver_vpc {  
    reads {  
        vpc.dstSw : exact;  
        vpc.customer : exact;  
        vpc.dstAddr : lpm;  
    }  
    actions {  
        pop_route_vpc;  
    }  
    size : 1024;  
}
```

Compute

10.0.0.4

10.0.0.4




10.0.0.5

ICMP answer to 10.0.2




## L2 – Ethernet (using P4 VPC)

Compute




Compute



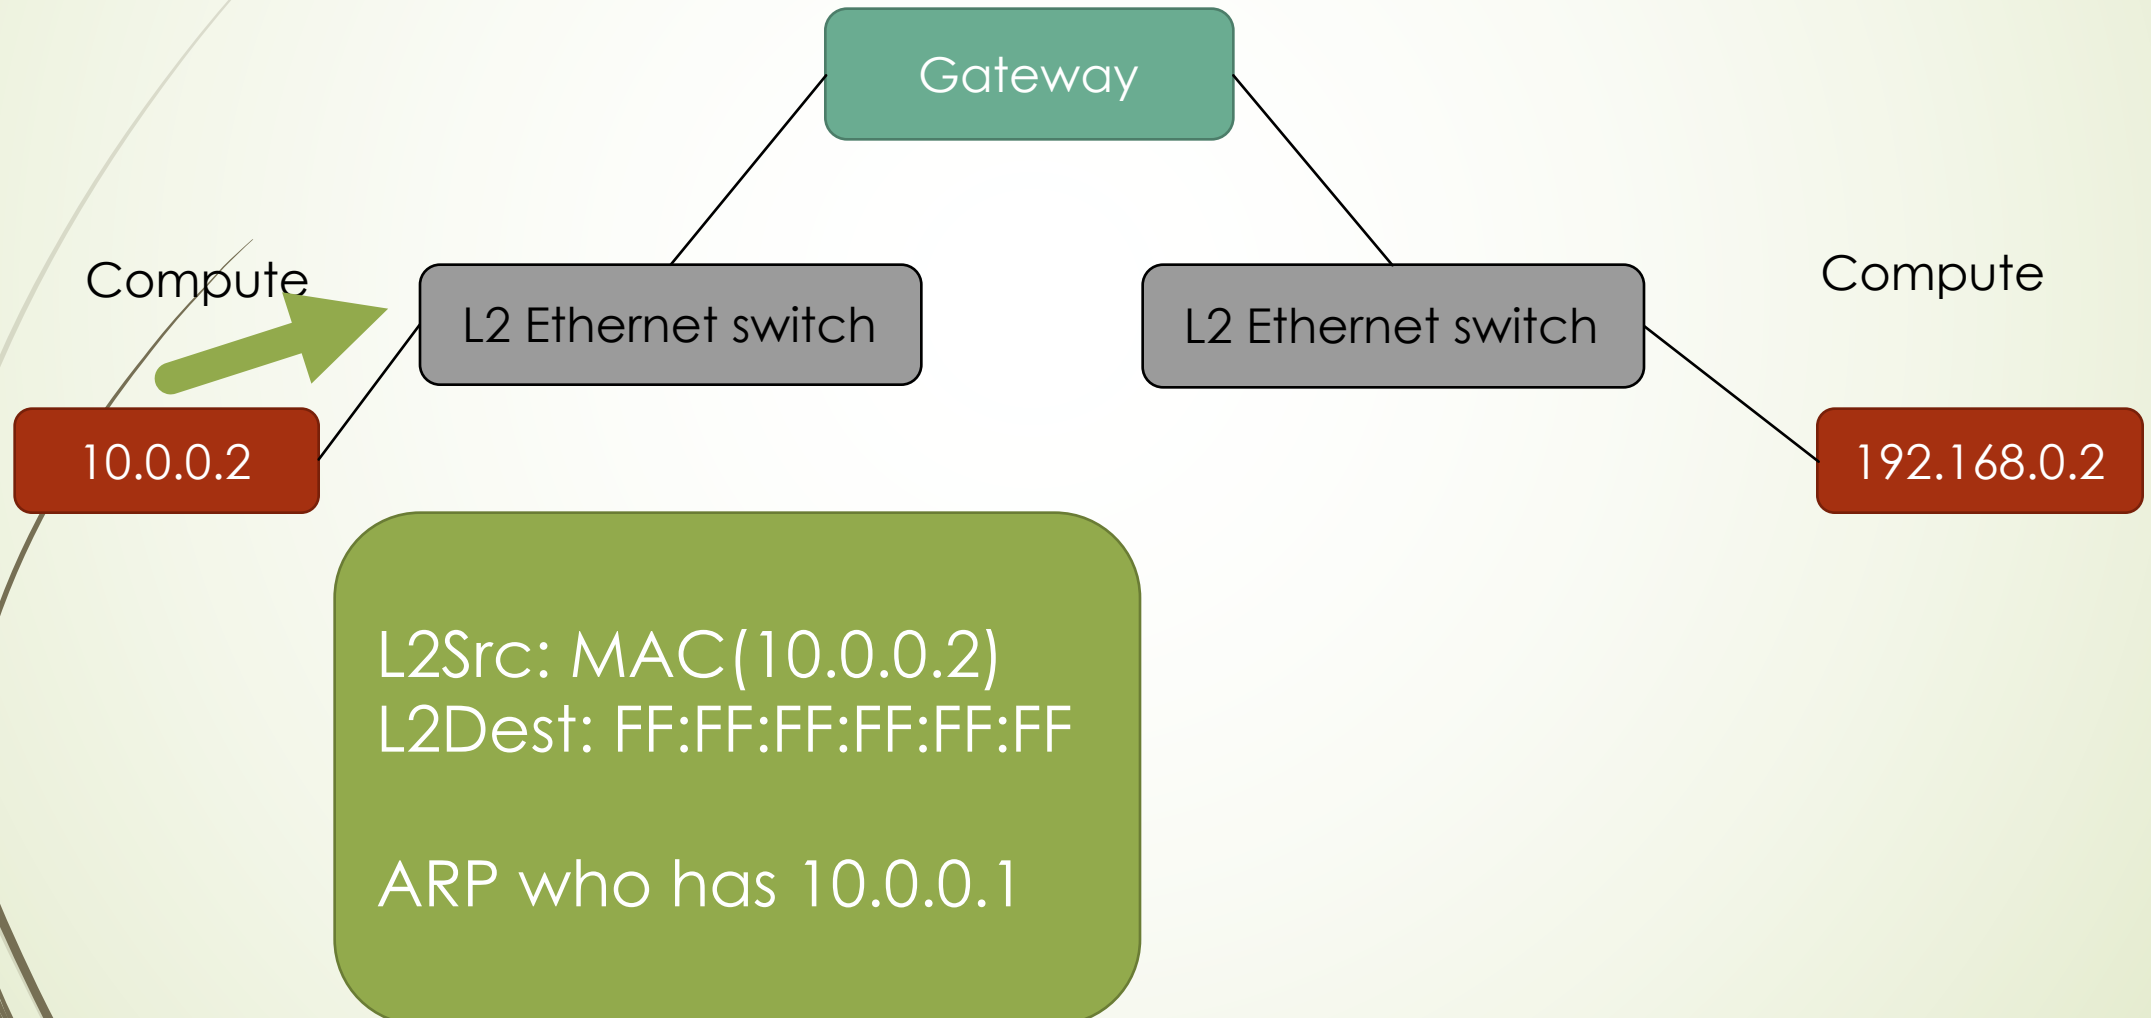
L2Src: MAC(10.0.0.2)  
L2Dest: MAC(10.0.0.4)

ICMP answer to 10.0.2

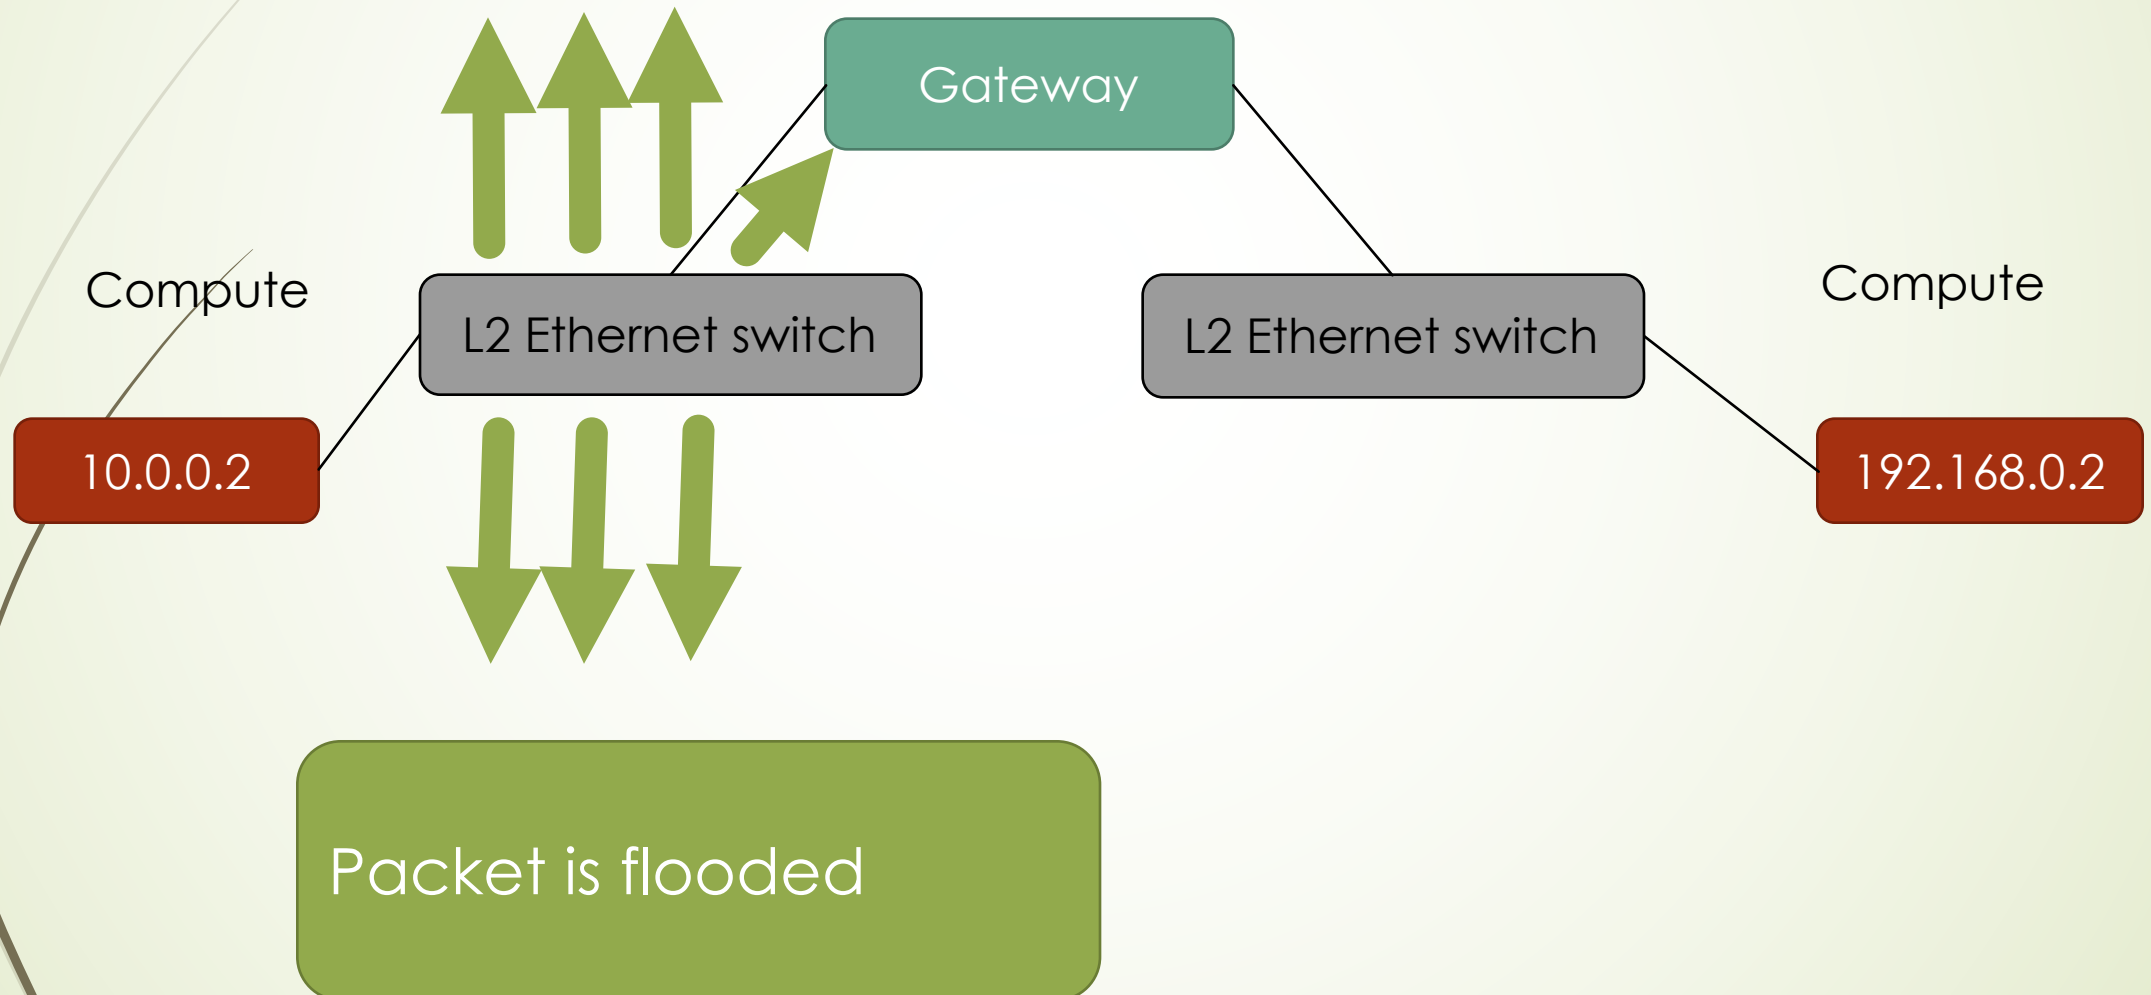
## L3 – IP Routing(standard)



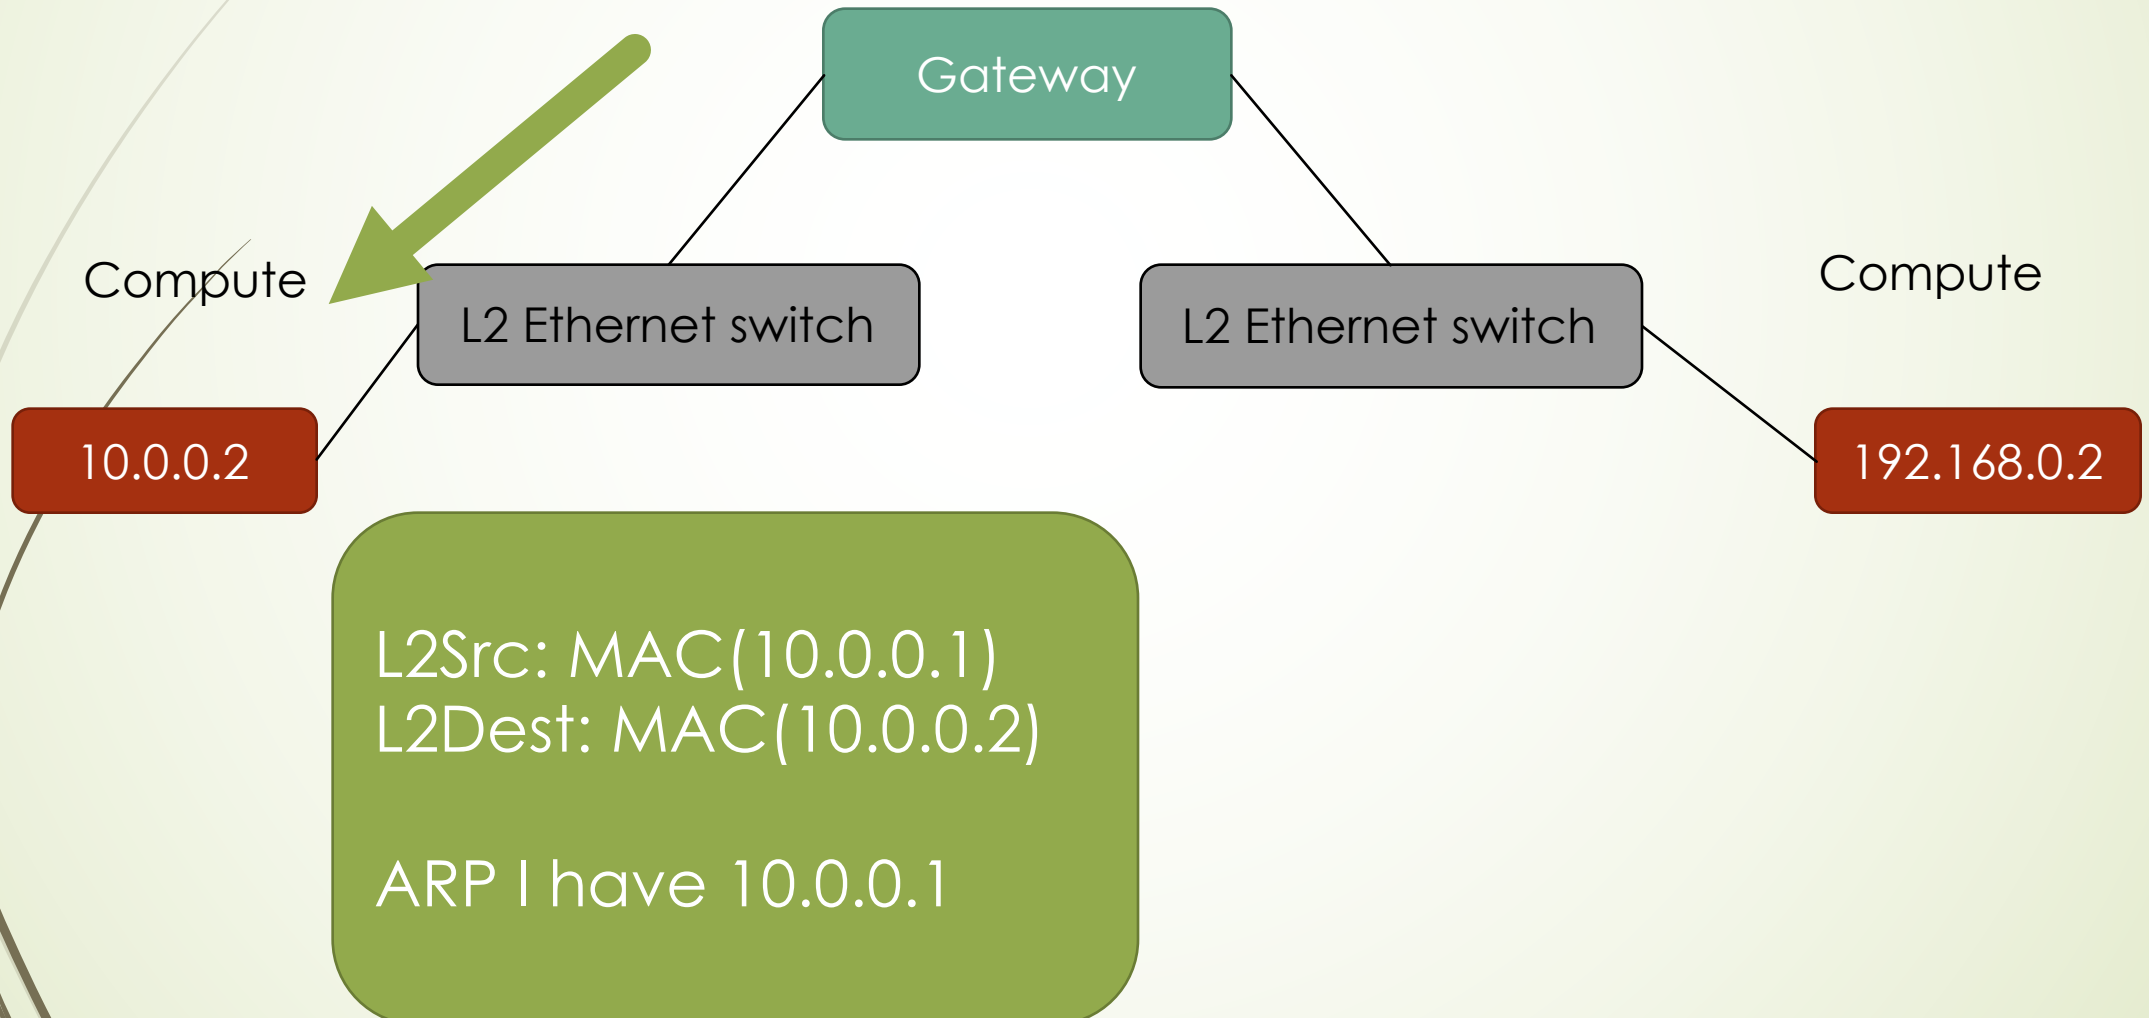
## L3 – IP Routing(standard)



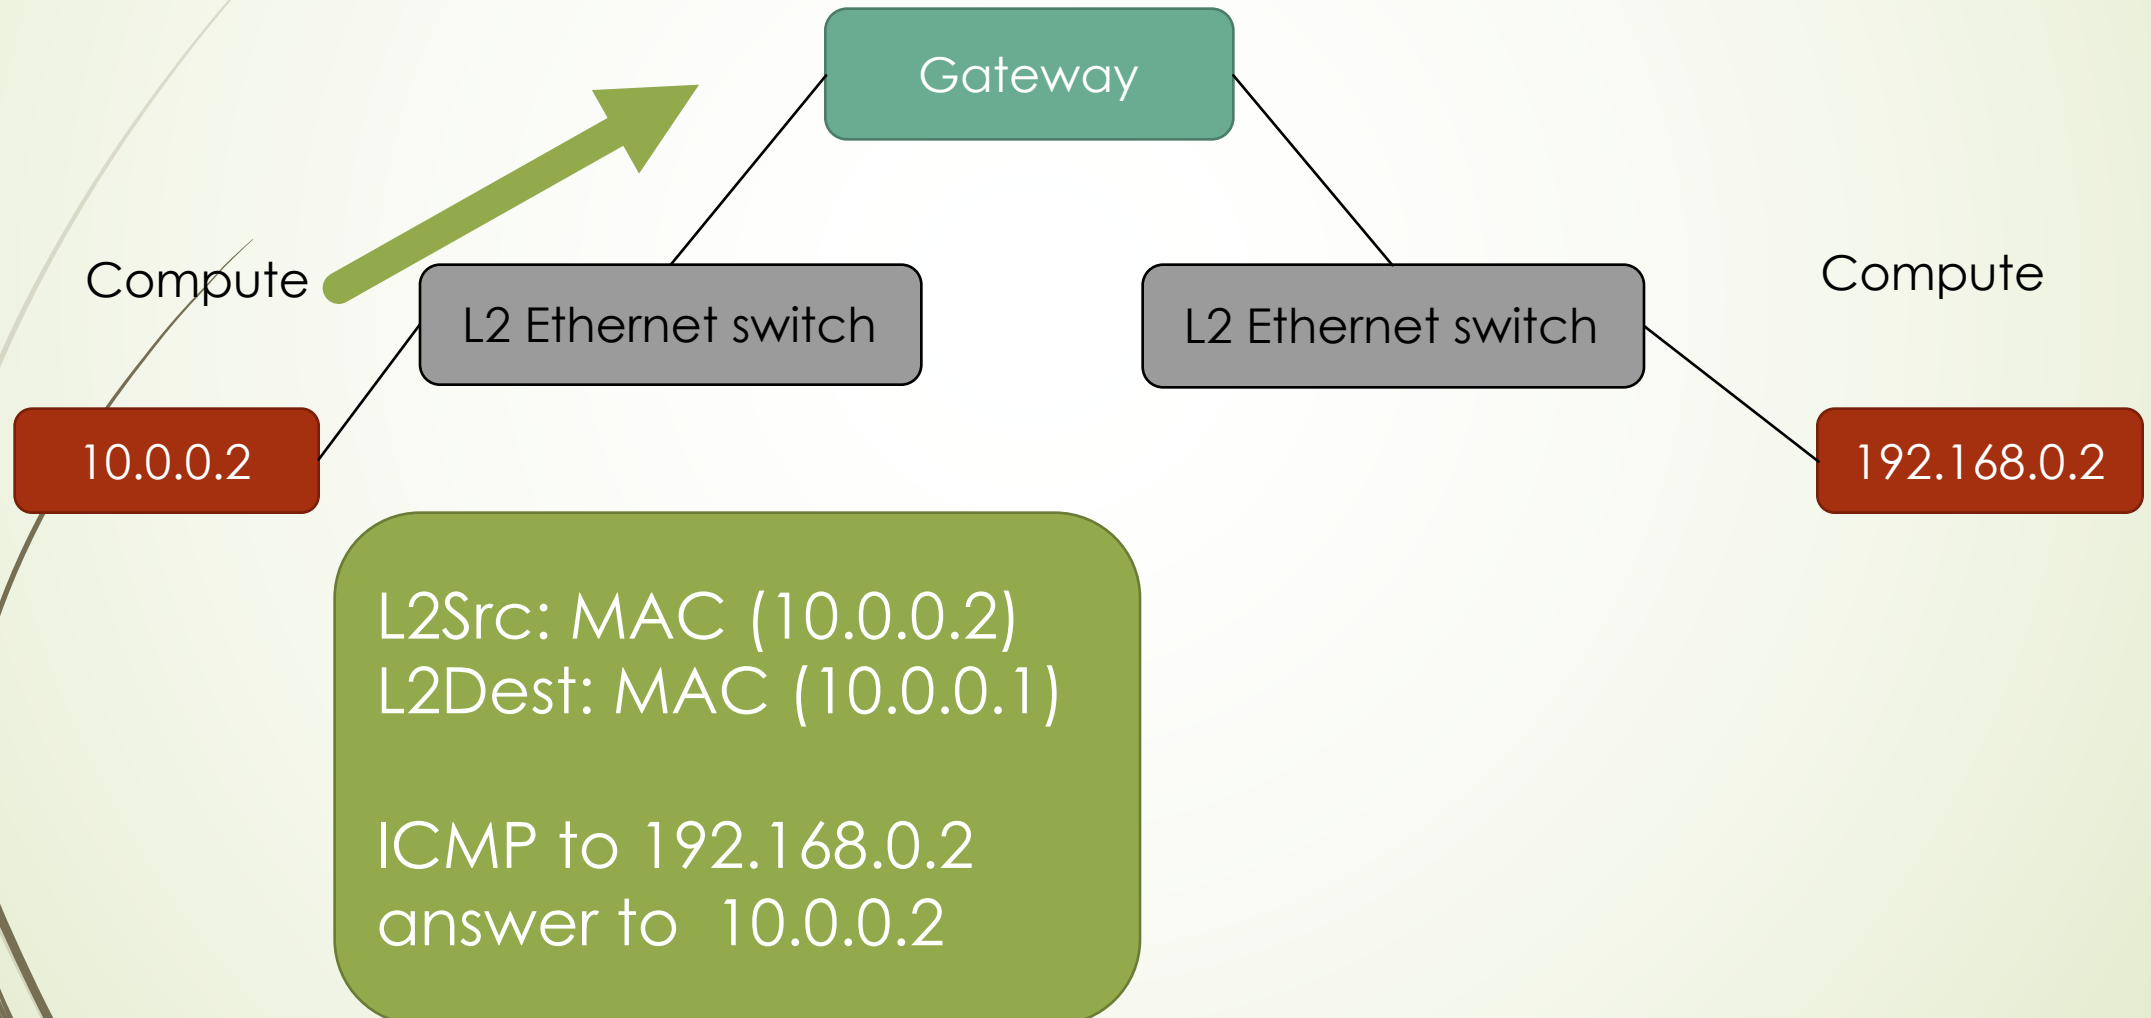
## L3 – IP Routing(standard)



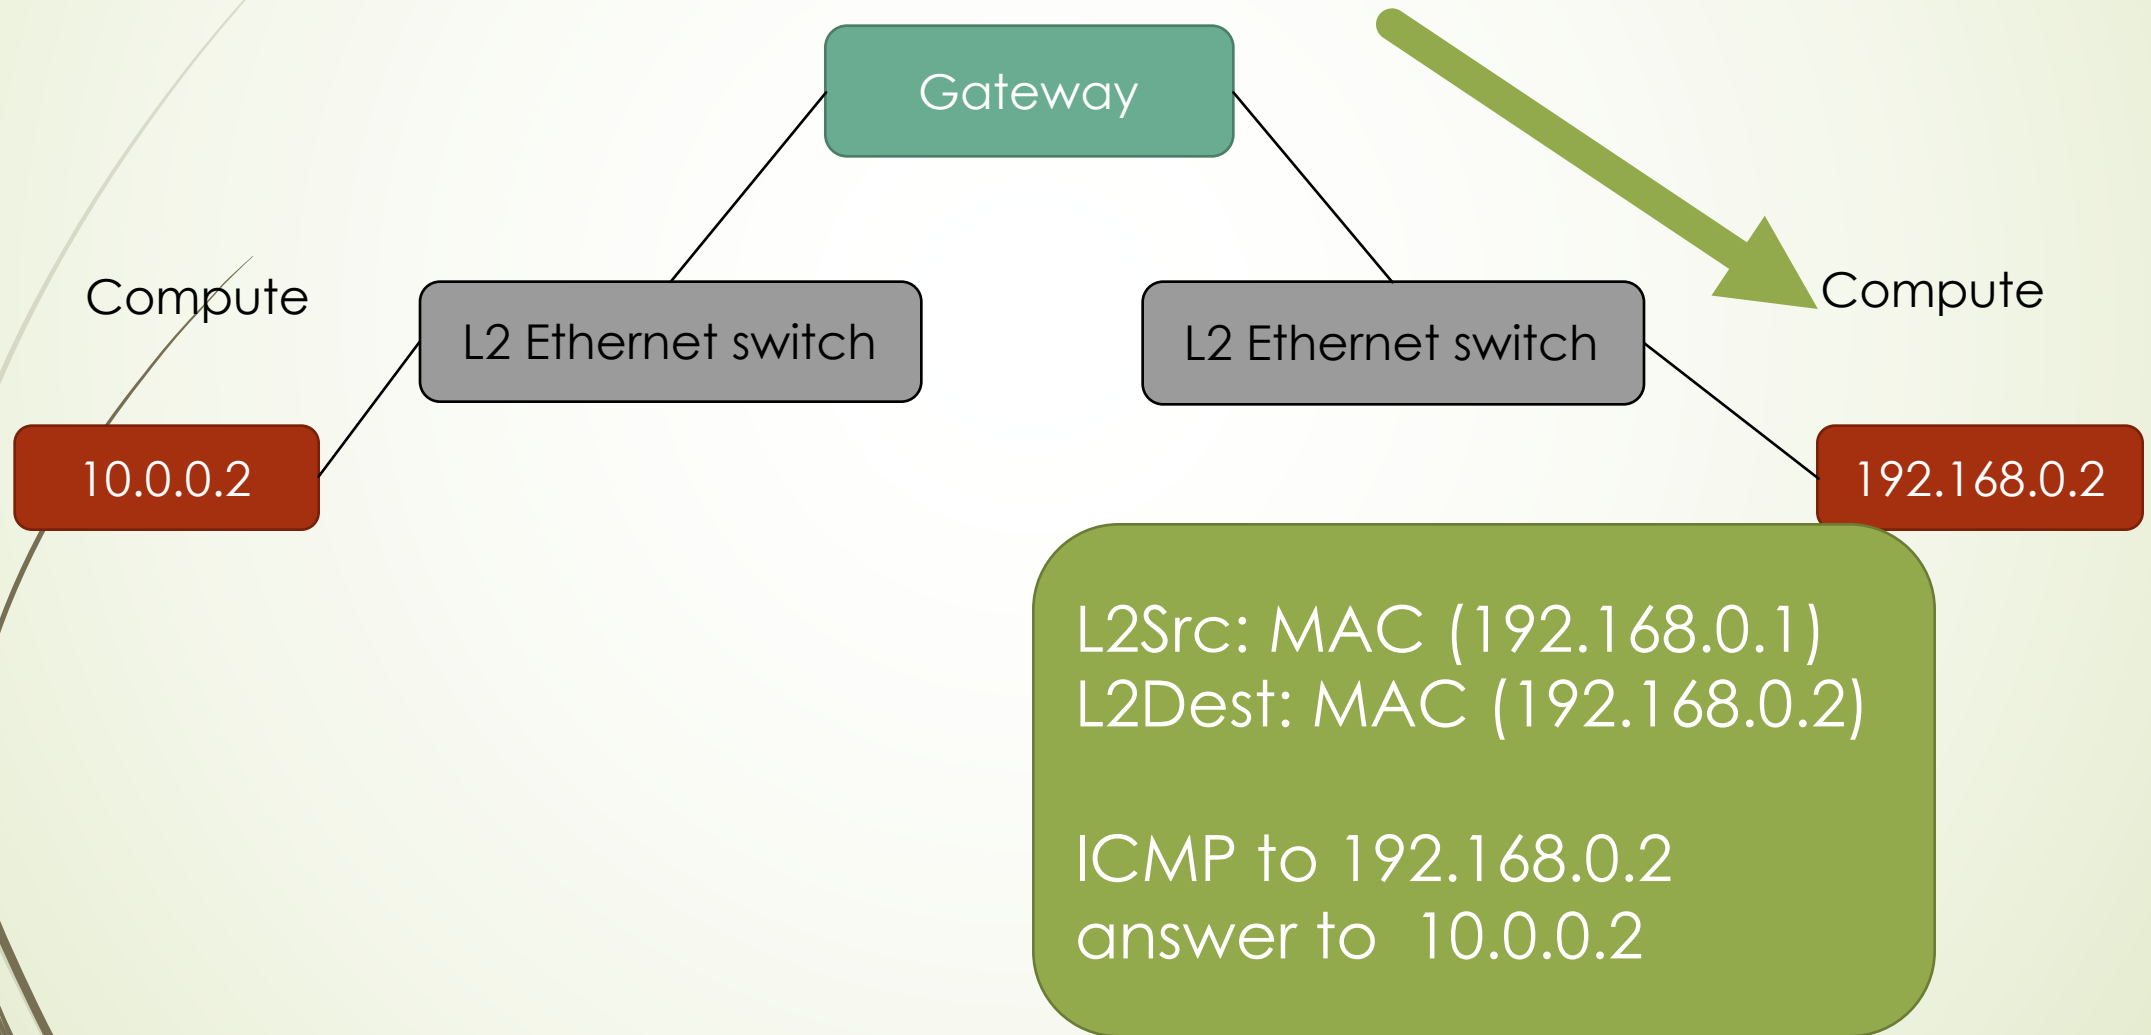
## L3 – IP Routing(standard)



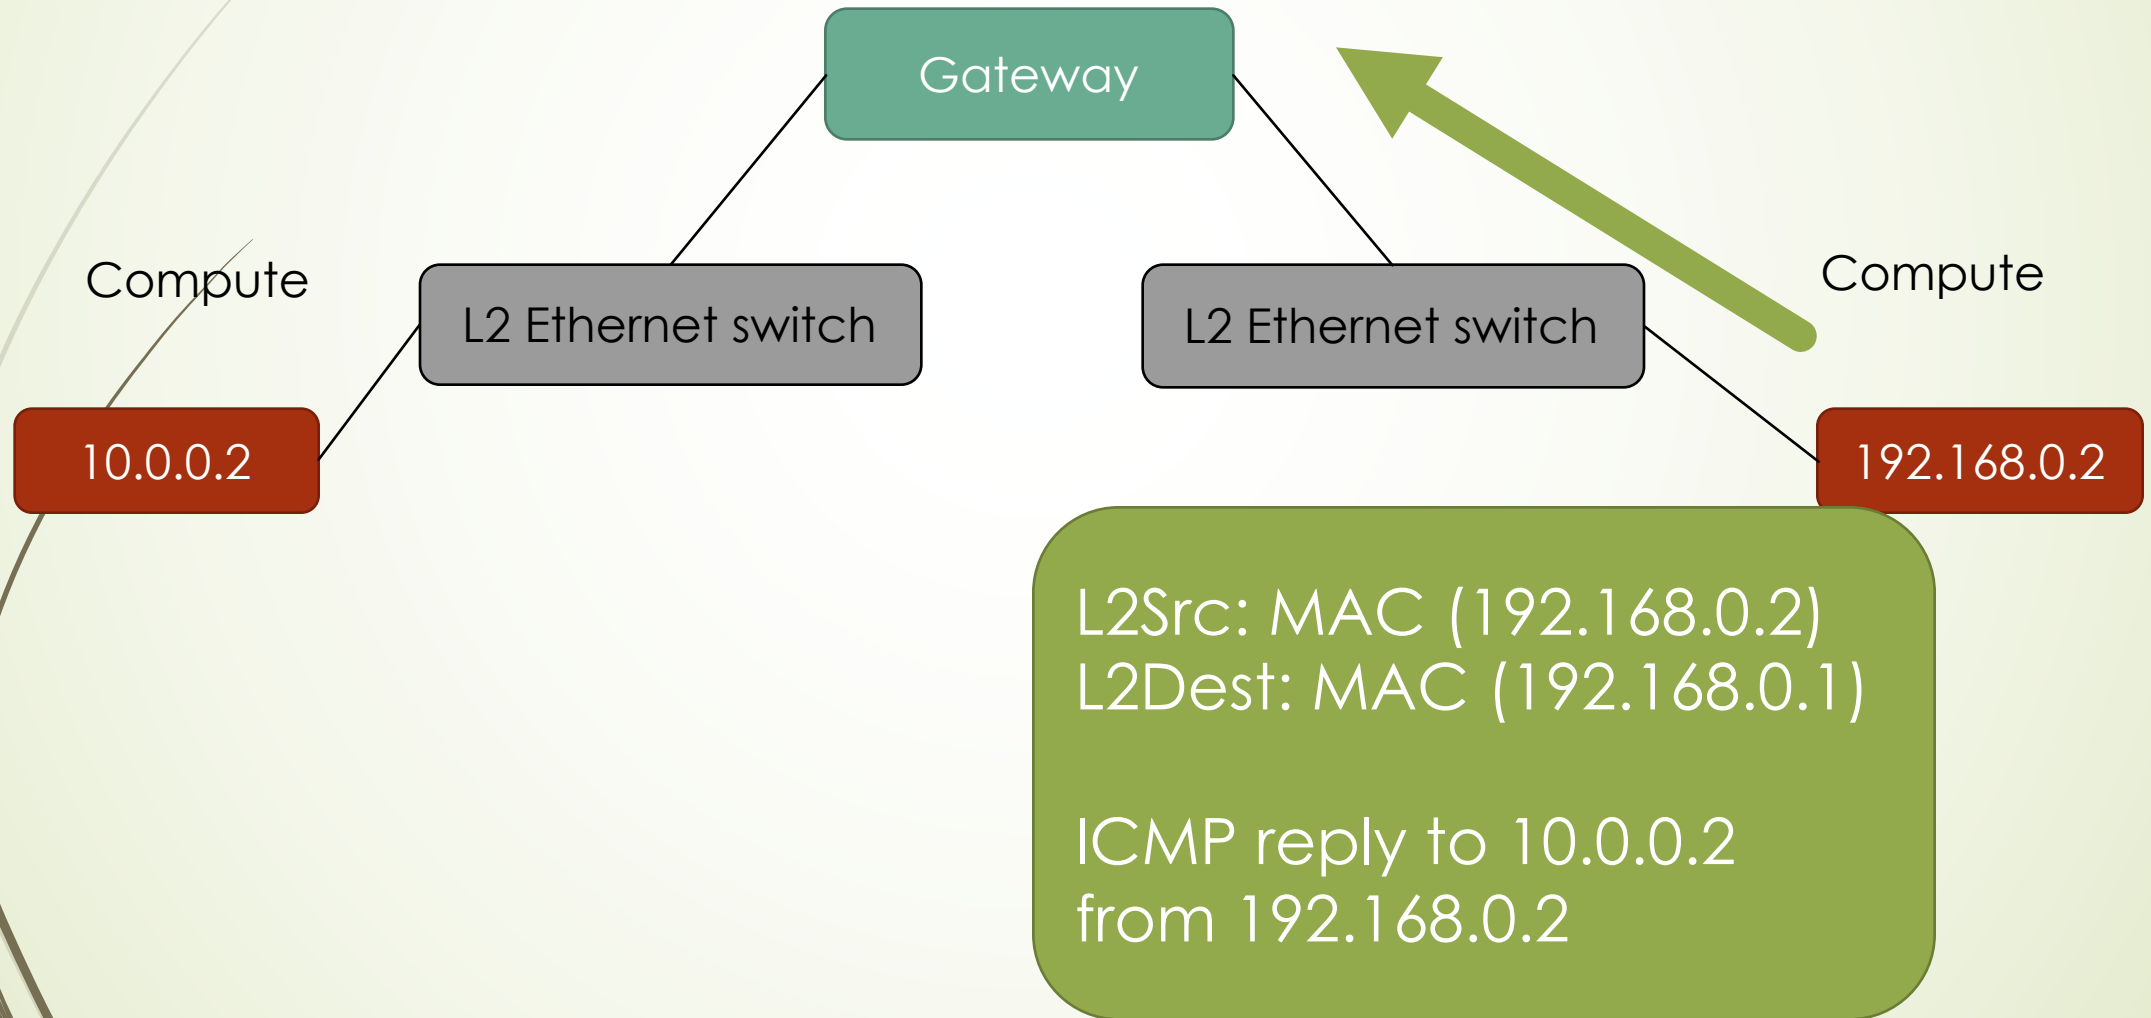
## L3 – IP Routing(standard)



## L3 – IP Routing(standard)

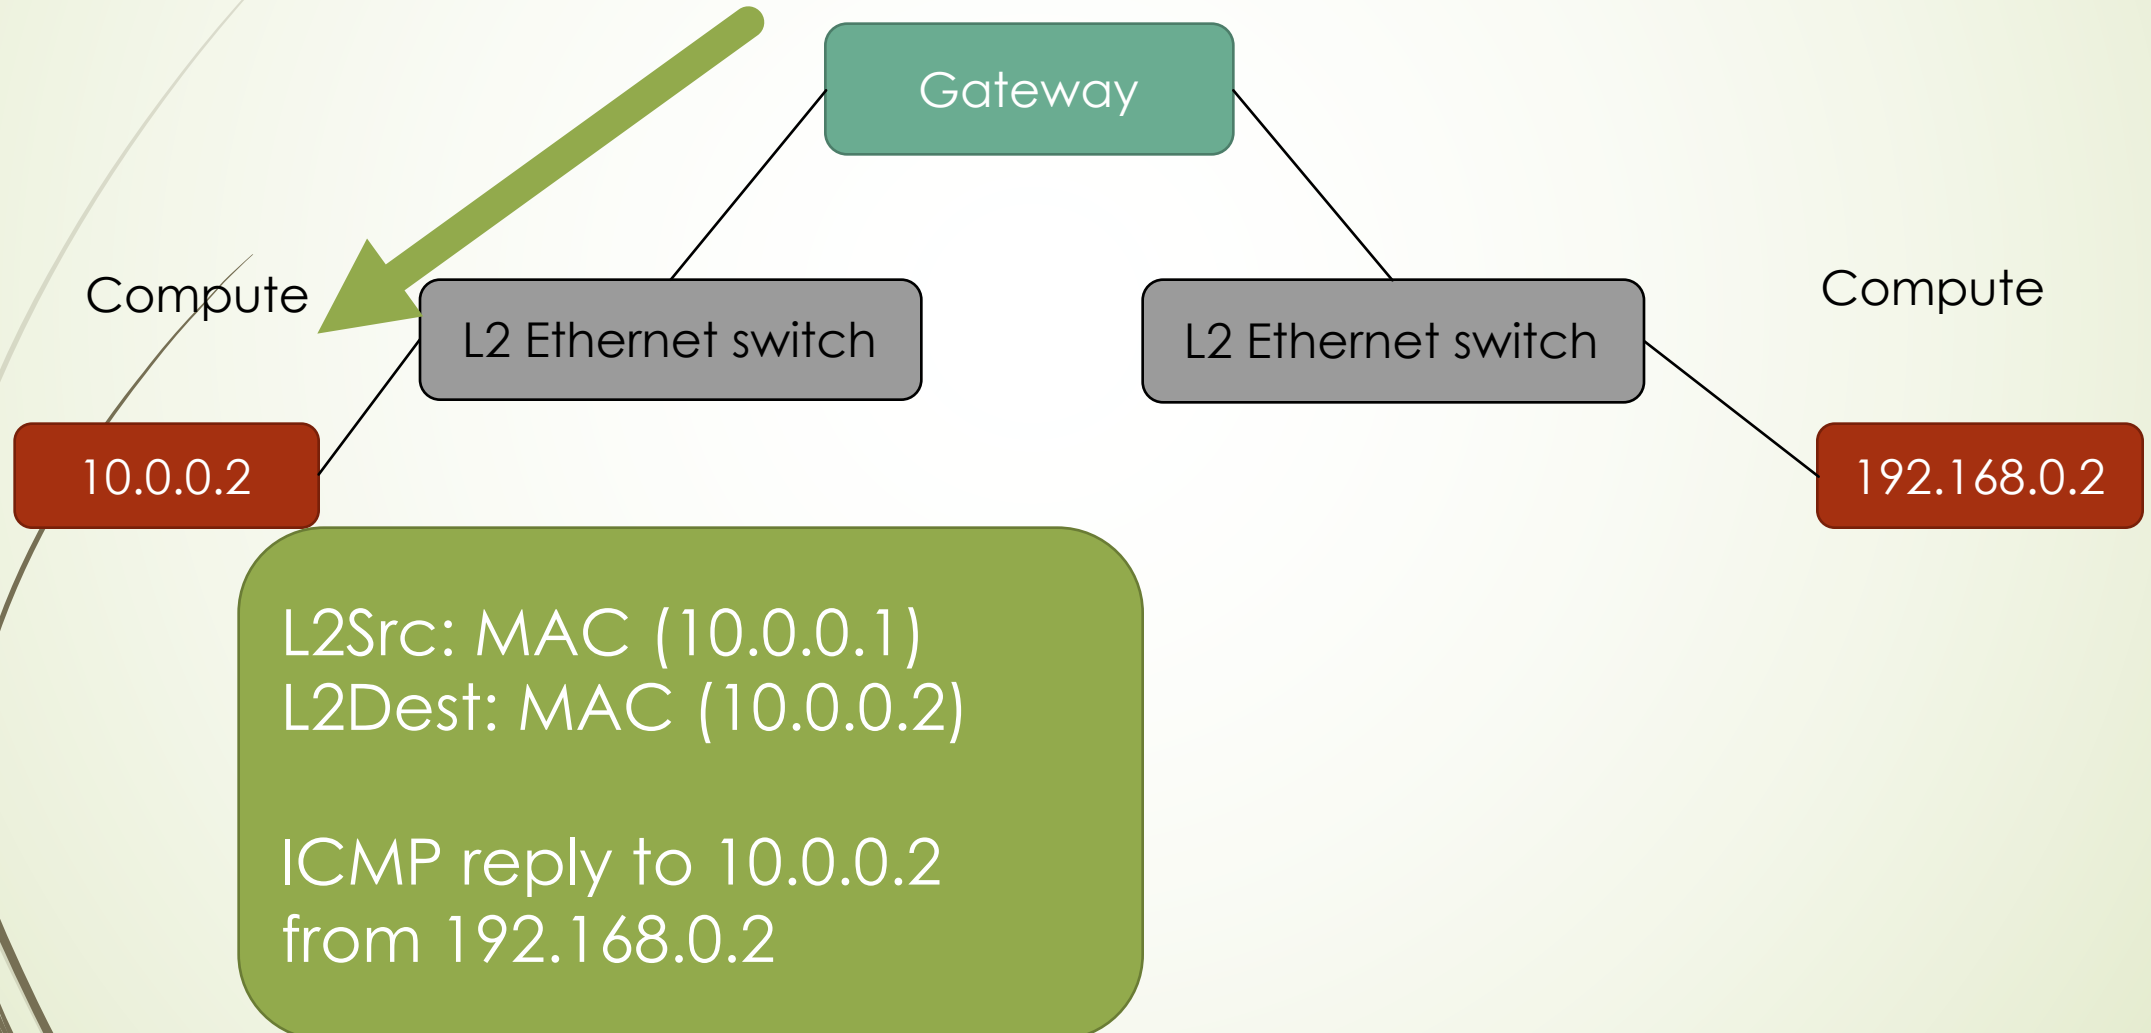


## L3 – IP Routing(standard)





## L3 – IP Routing(standard)



## L3 – IP Routing (using P4 VPC)

Compute

10.0.0.2

10.0.0.3

10.0.0.2

P4  
A

P4  
B


P4  
C

Compute

192.168.0.2

192.168.0.2

192.168.0.3



## L3 – IP Routing (using P4 VPC)

Compute

10.0.0.2

10.0.0.3

10.0.0.2



P4  
A

P4  
B

P4  
C

Compute

192.168.0.2

192.168.0.2

192.168.0.3

L2Src: MAC(10.0.0.2)  
L2Dest: FF:FF:FF:FF:FF:FF

ARP who has 10.0.0.1

## L3 – IP Routing (using P4 VPC)

Compute

10.0.0.2

Compute

192.168.0.2

**We do not have a GW in our network but we will reply a ARP reply with a fake GW MAC**

ARP who has 10.0.0.1

## L3 – IP Routing (using P4 VPC)

Compute

10.0.0.2

10.0.0.3

10.0.0.2

L2

L2

A

```
parser start {  
    return parse_ethernet;  
}  
  
parser parse_ethernet {  
    extract(ethernet);  
    return select(latest.etherType) {  
        ETHERTYPE_VPC : parse_vpc;  
        ETHERTYPE_IPV4 : parse_ipv4;  
        ETHERTYPE_ARP : parse_arp_rarp;  
        default: ingress;  
    }  
}
```

Compute

192.168.0.2

192.168.0.2

192.168.0.3

## L3 – IP Routing (using P4 VPC)

```
control ingress {
```

```
    if (ethernet.etherType == ETHERTYPE_ARP and arp_rarp.opcode == 1) {  
        apply(address_arp_packet);
```

```
    } else if (ethernet.etherType == ETHERTYPE_IPV4 ){  
        apply(address_ip_packet);  
    }
```

```
    if ((ethernet.etherType == ETHERTYPE_ARP and arp_rarp.opcode == 1) or  
        apply(vpc_customer);  
    }
```

```
    if (ethernet.etherType == ETHERTYPE_ARP and arp_rarp.opcode == 1) {  
        if (ingress_metadata.customer > 0){  
            apply(arp_reply);  
        }
```

```
    } else if (ethernet.etherType == ETHERTYPE_IPV4 ){  
        if (ingress_metadata.customer > 0){  
            apply(encapsulate_vpc);  
            apply(vpc_sw_id);  
            apply(vpc_dst);  
        }
```

compute

2.168.0.2

2.168.0.2

2.168.0.3

## L3 – IP Routing (using P4 VPC)

Compute

10.0.0.2

10.0.0.3

10.0.0.2

P4  
A

P4  
B

P4  
C

Compute

192.168.0.2

192.168.0.2

192.168.0.3

L2Src: MAC(10.0.0.1)  
L2Dest: MAC(10.0.0.2)

ARP I have 10.0.0.1



## L3 – IP Routing (using P4 VPC)

Compute

10.0.0.2

10.0.0.3

10.0.0.2



P4  
A

P4  
B

P4  
C

Compute

192.168.0.2

192.168.0.2

192.168.0.3

L2Src: MAC(10.0.0.2)  
L2Dest: MAC(10.0.0.1)

ICMP to 192.168.0.2  
answer to 10.0.2



## L3 – IP Routing (using P4 VPC)

Compute

10.0.0.2

10.0.0.3

10.0.0.2

```
parser start {  
    return parse_ethernet;  
}  
  
parser parse_ethernet {  
    extract(ethernet);  
    return select(latest.etherType) {  
        ETHERTYPE_VPC : parse_vpc;  
        ETHERTYPE_IPV4 : parse_ipv4;  
        ETHERTYPE_ARP : parse_arp_rarp;  
        default: ingress;  
    }  
}
```

answer to 10.0.2

Compute

192.168.0.2

192.168.0.2

192.168.0.3

## L3 – IP Routing (using P4 VPC)

Compute




P4  
A

P4  
B

P4  
C

Compute



Customer: Red

Src Switch: A

Dest Switch: C

Src IP: 10.0.0.2

Dest IP: 192.168.0.2

L2Src: MAC(10.0.0.2)




L2Dest: MAC(10.0.0.1)

ICMP to 192.168.0.2


answer to 10.0.2

## L3 – IP Routing (using P4 VPC)

Compute



Compute



Customer: Red

Src Switch: A

Dest Switch: C

Src IP: 10.0.0.2

Dest IP: 192.168.0.2

L2Src: MAC(10.0.0.2)

L2Dest: MAC(10.0.0.1)

ICMP to 192.168.0.2

answer to 10.0.2

## L3 – IP Routing (using P4 VPC)

Compute

10.0.0.2

10.0.0.3

10.0.0.2

```
}  
if (valid(vpc)){  
    apply(l2_addr);  
    apply(routing_pvc);  
    apply(deliiver_pvc);  
}  
}
```

Compute

192.168.0.2

192.168.0.2

192.168.0.3

10.0.0.2  
192.168.0.2

L2Src: MAC(10.0.0.2)  
L2Dest: MAC(10.0.0.1)  
ICMP to 192.168.0.2  
answer to 10.0.2

```
action set_l2_addr(srcAddr, dstAddr) {  
    modify_field(ethernet.srcAddr, srcAddr);  
    modify_field(ethernet.dstAddr, dstAddr);  
}
```

Computer

10.0.0.1

10.0.0.2

10.0.0.3

```
table l2_addr {  
    reads {  
        vpc.customer : exact;  
        vpc.dstSw : exact;  
        vpc.srcAddr : lpm;  
        vpc.dstAddr : exact;  
    }  
    actions {  
        _noop;  
        set_l2_addr;  
    }  
    size : 1024;  
}
```

Computer

192.168.0.2




192.168.0.2

168.0.3


answer to 10.0.2

## L3 – IP Routing (using P4 VPC)

Compute



Compute



Customer: Red

Src Switch: A

Dest Switch: C

Src IP: 10.0.0.2

Dest IP: 192.168.0.2

L2Src: MAC(192.168.0.1)

L2Dest: MAC(192.168.0.2)




ICMP 10 192.168.0.2

answer to 10.0.2




## L3 – IP Routing (using P4 VPC)

Compute



Compute



L2Src: MAC(192.168.0.1)  
L2Dest: MAC(192.168.0.2)  
ICMP to 192.168.0.2  
answer to 10.0.2





# Summary



- Custom parser that just covers ARP, IP and VPC headers
- Create a new encapsulation using 0x0777 Ethernet protocol VPC (custom and non standard with fields like customer id, src/destination IP Address, ...)
- Manipulation of packets:
  - Convert an ARP request into an ARP reply
  - Encapsulate IP packets with VPC (0x0777)
  - Route based on VPC header
  - Modify any field in the packet header
- Simulate a non existing GW with a custom protocol





# Demo



# Steps

- Ping from red customer to an ip address in the same subnet.
  - h102red ping 10.0.0.4
  - h102red nc 10.0.0.4 8888
  - h104red ifconfig eth0
- Ping from blue customer to same ip address as previous test
  - h102blue ping 10.0.0.4
  - h102blue nc 10.0.0.4 8888
  - h103blue ifconfig eth0
- Check same test with different ip subnet 192.168.0.3 h203blue or h203red

# Demo environment

Compute

10.0.0.2

10.0.0.3

10.0.0.2

10.0.0.3

P4  
A

P4  
B

P4  
C

Compute

10.0.0.4

10.0.0.5

192.168.0.2

192.168.0.3

10.0.0.4

10.0.0.5

192.168.0.2

192.168.0.3

# Demo environment

